

Lecture 12 - Dynamic Memory Management

Meng-Hsun Tsai
CSIE, NCKU

12.1 Dynamic Memory Allocation



Dynamic memory Allocation

- C's data structures, including arrays, are normally fixed in size.
- Fixed-size data structures can be a problem, since we're forced to choose their sizes before executing a program.
- Fortunately, C supports **dynamic memory allocation**: the ability to allocate memory during program execution.
- Using dynamic memory allocation, we can design data structures that grow (and shrink) as needed.

Dynamic memory Allocation (cont.)

- Dynamic memory allocation is **used most often for strings, arrays, and structures.**
- Dynamically allocated structures **can be linked together to form lists, trees, and other data structures.**
- Dynamic memory allocation **is done by calling a memory allocation function.**

Memory Allocation Functions

- The `<stdlib.h>` header declares three memory allocation functions:

`malloc`—Allocates a block of memory but **doesn't initialize it**.

`calloc`—Allocates a block of memory and **clears it**.

`realloc`—**Resizes** a previously allocated block of memory.

- **These functions return** a value of type `void *` (a “generic” pointer).

Null Pointers

- If a memory allocation function **can't locate** a memory block of the requested size, it **returns a *null pointer***.
- A **null pointer is a special value** that can be distinguished from all valid pointers.
- After we've **stored the function's return value** in a pointer variable, we must test **to see if it's a null pointer**.

Null Pointers (cont.)

- An example of testing `malloc`'s return value:

```
p = malloc(10000);  
if (p == NULL) {  
    /* allocation failed; take appropriate action */  
}
```

- `NULL` is a macro (defined in various library headers) that represents the null pointer.
- Some programmers combine the call of `malloc` with the `NULL` test:

```
if ((p = malloc(10000)) == NULL) {  
    /* allocation failed; take appropriate action */  
}
```

Null Pointers (cont.)

- Pointers test true or false in the same way as numbers.
- All non-null pointers test true; only null pointers are false.

- Instead of writing

```
if (p == NULL) ...  
we could write
```

```
if (!p) ...
```

- Instead of writing

```
if (p != NULL) ...  
we could write
```

```
if (p) ...
```


12.2 Dynamically Allocated Strings



Dynamically Allocated Strings

- Dynamic memory allocation is often useful for working with strings.
- **Strings** are stored in character arrays, and it **can be hard to anticipate how long these arrays need to be.**
- **By allocating strings dynamically, we can postpone the decision until the program is running.**

Using `malloc` to Allocate Memory for a String

- Prototype for the `malloc` function:

```
void *malloc(size_t size);
```

- `malloc` **allocates** a block of **size bytes** and **returns a pointer** to it.
- `size_t` is an **unsigned integer** type defined in the library.

Using `malloc` to Allocate Memory for a String (cont.)

- A call of `malloc` that allocates memory for a string of `n` characters:

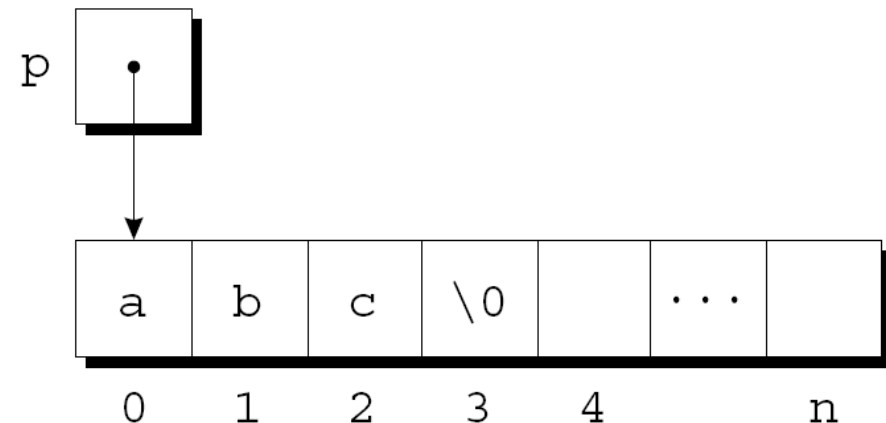
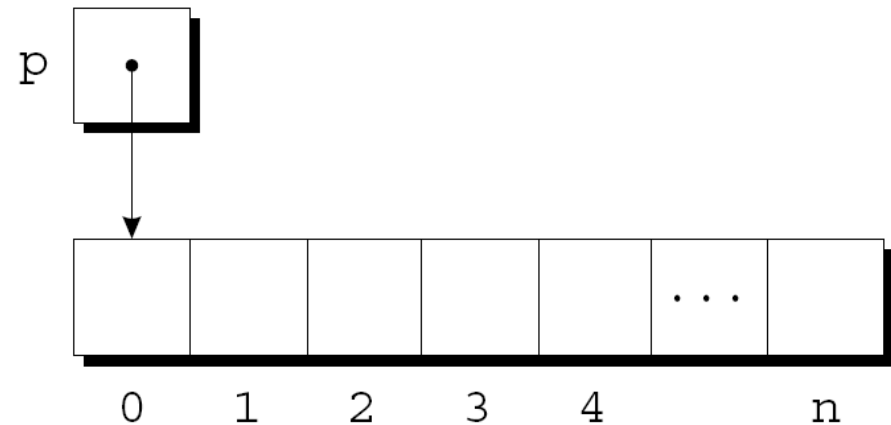
```
char * p;  
p = malloc(n + 1);
```

- Each character requires one byte of memory; adding 1 to `n` leaves room for the null character.
- Some programmers prefer to cast `malloc`'s return value, although the cast is not required:

```
p = (char *) malloc(n + 1);
```

Using `malloc` to Allocate Memory for a String (cont.)

- Memory allocated using `malloc` **isn't cleared**, so `p` will point to an **uninitialized array of $n + 1$ characters**:
- **Calling `strcpy` is one way to initialize** this array:
`strcpy(p, "abc");`
- The first four characters in the array will now be `a`, `b`, `c`, and `\0`:



Using Dynamic memory Allocation in String Functions

- Dynamic memory allocation makes it possible to write functions that return a pointer to a “new” string.
- Consider the problem of **writing a function that concatenates two strings without changing either one.**
- The function will **measure the lengths of the two strings** to be concatenated, then **call `malloc` to allocate the right amount of space** for the result.

Using Dynamic memory Allocation in String Functions (cont.)

```
char *concat(const char *s1, const char *s2)
{
    char *result;

    result = malloc(strlen(s1) + strlen(s2) + 1);
    if (result == NULL) {
        printf("Error: malloc failed in concat\n");
        exit(EXIT_FAILURE);
    }
    strcpy(result, s1);
    strcat(result, s2);
    return result;
}
```

Using Dynamic memory Allocation in String Functions (cont.)

- A call of the `concat` function:

```
p = concat("abc", "def");
```

- After the call, **p will point to the string "abcdef"**, which is stored in a dynamically allocated array.

Using Dynamic memory Allocation in String Functions (cont.)

- Functions such as `concat` that dynamically allocate memory **must be used with care**.
- **When the string that `concat` returns is no longer needed, we'll want to call the `free` function to release the space that the string occupies.**
- **If we don't, the program may eventually run out of memory.**

Program: Printing a One-Month Reminder List (Revisited)

- The `remind2.c` program is based on the `remind.c` program of Lecture 10, which **prints a one-month list of daily reminders**.
- The **original `remind.c` program stores reminder strings in a two-dimensional array of characters**.
- **In the new program, the array will be one-dimensional; its elements will be pointers to dynamically allocated strings**.

Program: Printing a One-Month Reminder List (Revisited) (cont.)

- **Advantages** of switching to dynamically allocated strings:
 - Uses space more efficiently by allocating the exact number of characters needed to store a reminder.
 - Avoids calling `strcpy` to move existing reminder strings in order to make room for a new reminder.
- Switching from a two-dimensional array to an array of pointers requires **changing only eight lines** of the program (shown in **bold**).

Program: Printing a One-Month Reminder List (Revisited) (cont.)

remind2.c

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

#define MAX_REMIND 50    /* maximum number of reminders */
#define MSG_LEN 60      /* max length of reminder message */

int read_line(char str[], int n);
int main(void)
{
    char *reminders[MAX_REMIND];
    char day_str[3], msg_str[MSG_LEN+1];
    int day, i, j, num_remind = 0;
```

```
char reminders[MAX_REMIND][MSG_LEN+3];
```

Program: Printing a One-Month Reminder List (Revisited) (cont.)

```
for (;;) {
    if (num_remind == MAX_REMIND) {
        printf("-- No space left --\n");
        break;
    }
    printf("Enter day and reminder: ");
    scanf("%2d", &day);
    if (day == 0)
        break;
    sprintf(day_str, "%2d", day);
    read_line(msg_str, MSG_LEN);

    for (i = 0; i < num_remind; i++)
        if (strcmp(day_str, reminders[i]) < 0)
            break;
    for (j = num_remind; j > i; j--)
        reminders[j] = reminders[j-1];
```

```
strcpy(reminders[j],
       reminders[j-1]);
```

Program: Printing a One-Month Reminder List (Revisited) (cont.)

```
reminders[i] = malloc(2 + strlen(msg_str) + 1);  
if (reminders[i] == NULL) {  
    printf("-- No space left --\n");  
    break;  
}
```

```
strcpy(reminders[i], day_str);  
strcat(reminders[i], msg_str);
```

```
num_remind++;  
}
```

```
printf("\nDay Reminder\n");  
for (i = 0; i < num_remind; i++)  
    printf(" %s\n", reminders[i]);
```

```
return 0;
```

12.3 Dynamically Allocated Arrays



Dynamically Allocated Arrays

- The **close relationship between arrays and pointers** makes a **dynamically allocated array** as **easy to use** as an ordinary array.
- Suppose a program **needs an array of n integers**, where n is computed during program execution.
- We'll **first declare a pointer variable**:

```
int *a;
```
- **Once the value of n is known**, the program can **call `malloc` to allocate space** for the array:

```
a = malloc(n * sizeof(int));
```
- **Always use the `sizeof` operator** to calculate the amount of space required for each element.

Using `malloc` to Allocate memory for an Array (cont.)

- We can now ignore the fact that `a` is a pointer and use it instead as an array name, thanks to the relationship between arrays and pointers in C.
- For example, we could use the following loop to initialize the array that `a` points to:

```
for (i = 0; i < n; i++)  
    a[i] = 0;
```
- We also have the option of using pointer arithmetic instead of subscripting to access the elements of the array.

The `calloc` Function

- The `calloc` function is an **alternative to** `malloc`.
- Prototype for `calloc`:

```
void *calloc(size_t nmemb, size_t size);
```
- Properties of `calloc`:
 - Allocates space for an array with `nmemb` **elements**, **each of which** is `size` **bytes long**.
 - **Returns a null pointer if** the requested **space isn't available**.
 - **Initializes** allocated memory by setting **all bits to 0**.

The `calloc` Function (cont.)

- A call of `calloc` that allocates space for **an array of `n` integers**:

```
a = calloc(n, sizeof(int));
```

- By calling `calloc` with 1 as its first argument, we can **allocate space for a data item of any type**:

```
struct point { int x, y; } *p;
```

```
p = calloc(1, sizeof(struct point));
```

The `realloc` Function

- The `realloc` function can resize a dynamically allocated array.
- Prototype for `realloc`:

```
void *realloc(void *ptr, size_t size);
```
- `ptr` must point to a memory block obtained by a previous call of `malloc`, `calloc`, or `realloc`.
- `size` represents the new size of the block, which may be larger or smaller than the original size.

```
q = realloc(p, 20000);
```

The `realloc` Function (cont.)

- Properties of `realloc`:
 - When it expands a memory block, `realloc` doesn't initialize the bytes that are added to the block.
 - If `realloc` can't enlarge the memory block as requested, it returns a null pointer; the data in the old memory block is unchanged.
 - If `realloc` is called with a null pointer as its first argument, it behaves like `malloc`.
 - If `realloc` is called with 0 as its second argument, it frees the memory block.

The `realloc` Function (cont.)

- We expect `realloc` to be reasonably efficient:
 - When asked to reduce the size of a memory block, `realloc` should shrink the block “in place.”
 - `realloc` should always attempt to expand a memory block without moving it.
- If it can't enlarge a block, `realloc` will allocate a new block elsewhere, then copy the contents of the old block into the new one.
- Once `realloc` has returned, be sure to update all pointers to the memory block in case it has been moved.

12.4 Deallocating Memory




Deallocating memory

- `malloc` and the other memory allocation functions obtain memory blocks from a memory pool known as the **heap**.
- Calling these functions too often—or asking them for large blocks of memory—can exhaust the heap, causing the functions to return a null pointer.
- To make matters worse, a program may allocate blocks of memory and then lose track of them, thereby wasting space.

Deallocating memory (cont.)

```
p = malloc(...);  
q = malloc(...);  
p = q;
```

- A snapshot **after the first two statements** have been executed:
- **After q is assigned to p , both variables now point to the second memory block.**
- There are **no pointers to the first block**, so we'll never be able to use it again.



Deallocating memory (cont.)

- A block of memory that's **no longer accessible** to a program is **said to be *garbage***.
- A program that **leaves garbage behind** has a ***memory leak***.
- Some languages provide a ***garbage collector*** that automatically locates and recycles garbage, **but C doesn't**.
- Instead, **each C program is responsible for recycling its own garbage by calling the `free` function** to release unneeded memory.

The **free** Function

- Prototype for `free`:

```
void free(void *ptr);
```

- `free` will be passed a pointer to an unneeded memory block:

```
p = malloc(...);
```

```
q = malloc(...);
```

```
free(p);
```

```
p = q;
```

- Calling `free` releases the block of memory that `p` points to.

The “Dangling Pointer” Problem

- Using `free` leads to a new problem: ***dangling pointers***.
- `free(p)` deallocates the memory block that `p` points to, but doesn't change `p` itself.
- If we forget that `p` no longer points to a valid memory block, chaos may ensue:

```
char *p = malloc(4);  
...  
free(p);  
...  
strcpy(p, "abc");    /** WRONG **/
```

- Modifying the memory that `p` points to is a **serious error**.

The “Dangling Pointer” Problem (cont.)

- Dangling pointers can be hard to spot, since several pointers may point to the same block of memory.
- When the block is freed, all the pointers are left dangling.

A Quick Review to This Lecture

- The `<stdlib.h>` header declares three memory allocation functions which return **void pointers** (`void *`) to allocated space or **return null pointers** if allocation is failed:

- `malloc` **allocates** a block of **size bytes** but **doesn't initialize it**.

```
void *malloc(size_t size);
```

```
p = malloc(10000);  
if (p == NULL) { /* failed*/ }
```

- `calloc` allocates an array with `nmemb` **elements**, **each of which** is **size bytes long**, and **clears it**.

```
a = calloc(n, sizeof(int));
```

```
void *calloc(size_t nmemb, size_t size);
```

- `realloc` **resizes** a block (where `ptr` points to) to a new size `size`

```
void *realloc(void *ptr, size_t size);
```

A Quick Review to This Lecture (cont.)

- Each C program is responsible for recycling its own garbage by calling the `free` function to release unneeded memory.

```
void free(void *ptr);
```

```
p = malloc(...);  
free(p);
```

- When the block is freed, all the pointers are left dangling.